

Andrey Krisanov

Moscow, Russia | develop4net@gmail.com | github.com/akrisanov | linkedin.com/in/akrisanov | akrisanov.com

Staff Software Engineer | AI Infrastructure & ML Systems

PROFILE

Staff Software Engineer building AI infrastructure, ML systems, and distributed systems.

At Severstal, I lead the architecture and hands-on development of DaVinci, a shared GenAI platform running production workloads across multiple data centers on Kubernetes and NVIDIA GPUs. My work covers platform architecture, GPU infrastructure and capacity planning, reliability, observability, model lifecycle, and production LLM serving.

Previously, I built and scaled backend and platform systems across SaaS, fintech, data privacy, payments, and high-traffic consumer products.

SELECTED IMPACT

- Lead the architecture and hands-on development of Severstal’s shared GenAI platform across Kubernetes, NVIDIA GPU infrastructure, model serving, traffic management, observability, and production reliability; designed its target topology, capacity allocation, and resilience model for expansion to 48 H200, 8 H100, and 24 B200 GPUs across two data centers.
- Modernized a \$3M+ ARR SaaS platform, improving critical backend paths by 2–10x, achieving 99.998% availability, and reducing mean time to recovery by 50%.
- Designed and launched X5 Media, a platform that reached more than 20 million monthly active users within its first year.

WORK EXPERIENCE

Staff Software Engineer
Severstal | DaVinci GenAI Platform

Jan 2026 — Present
Moscow, Russia

- Lead architecture and technical direction for DaVinci, a shared GenAI platform on top of Kubernetes and NVIDIA GPUs supporting enterprise AI products, assistants, and coding agents.
- Own production LLM serving for self-hosted open-weight models using vLLM, including model placement, parallelism, batching, context limits, request prioritization, and resource allocation.
- Design and develop LLM-aware routing and inference control plane with controlled model rollouts and rollbacks.
- Benchmark and tune serving configurations against production traffic, balancing TTFT, throughput, response quality, and serving cost.
- Designed capacity allocation, degradation policies, failover, and recovery for expansion to 48 H200, 8 H100, and 24 B200 GPUs across two data centers.
- Build SLO-driven operations with Prometheus, Grafana, GPU telemetry, alerting, load testing, incident analysis, and capacity forecasting.

Chief Technology Officer
PlanFact | B2B Accounting SaaS

Jan 2024 — Dec 2025
Moscow, Russia

- Led the technical modernization of a \$3M+ ARR SaaS product in a 100-person company while managing a 25-person technology organization.
- Modernized the core .NET platform through a domain-oriented modular architecture, replacing legacy dependencies without disrupting business-critical services.
- Improved critical backend paths by 2–10x and scaled read-heavy workloads through targeted refactoring, SQL and database tuning, strongly consistent replication, and horizontal read scaling.
- Introduced service-level observability and incident-management practices, achieving 99.998% availability and reducing mean time to recovery by 50%.

Senior Software Engineer, Backend & Platform

Stattice.ai | Data Privacy Platform

Jun 2021 — Dec 2023

Berlin, Germany

- Built Python, FastAPI, and Celery backend services and Kubernetes platform workflows on GCP for a cloud PII-protection product and privacy-preserving machine-learning workloads.
- Modernized SDK architecture and cross-platform x64/ARM build pipelines; introduced type checking, static analysis, CI improvements, and security scanning while maintaining 99.99% availability.

Engineering Lead, Backend & Platform

X5 | Retail Technology

Jun 2019 — Jun 2021

Moscow, Russia

- Designed and launched X5 Media, a Python, FastAPI, and Next.js platform that reached more than 20 million monthly active users within its first year.
- Owned backend and platform architecture during the transition from outsourced to in-house engineering; built the first cross-functional teams and halved time-to-market for existing products.

Earlier Engineering Experience

Fintech, SaaS, and Startups

2012 — 2019

Germany, Russia, and Remote

- **Independent consultant:** Migrated a Frankfurt fintech from on-premises infrastructure to AWS, cutting infrastructure costs by 50%; designed cloud architectures for two early-stage startups.
- **Clark:** Built a secure server-to-server integration with N26 that enabled in-app insurance signing and created a new partnership revenue stream; modernized and hardened a Ruby on Rails platform to improve security, performance, maintainability, and support for white-label products.
- **Ridewithlocal:** Owned backend engineering, AWS infrastructure, and production reliability for the initial product launch, maintaining 99.99% uptime; delivered the MVP on schedule and helped the company secure a \$300K pre-seed round.
- **Payler:** Designed and launched a PCI DSS-certified payment gateway, helping transform the company from a services business into a product-focused fintech platform; scaled engineering from one developer to a team of approximately 30.

SKILLS

- **Systems:** Distributed Systems, LLM Model Serving, AI Gateways, GPU Infrastructure, NVIDIA GPUs
- **Programming Languages:** Go, Python, C# (.NET), TypeScript, Ruby
- **Databases:** PostgreSQL, ClickHouse, Redis
- **Infrastructure:** Linux, Docker, Kubernetes, Ansible, Terraform, AWS, GCP, Yandex Cloud
- **Observability:** Prometheus, Grafana, DCGM, OpenTelemetry
- **Leadership:** Technical Leadership, Engineering Management, Team Building & Mentoring

EDUCATION

Bryansk State University

Master's in Applied Mathematics & Information Science

Bryansk, Russia

2006 — 2011